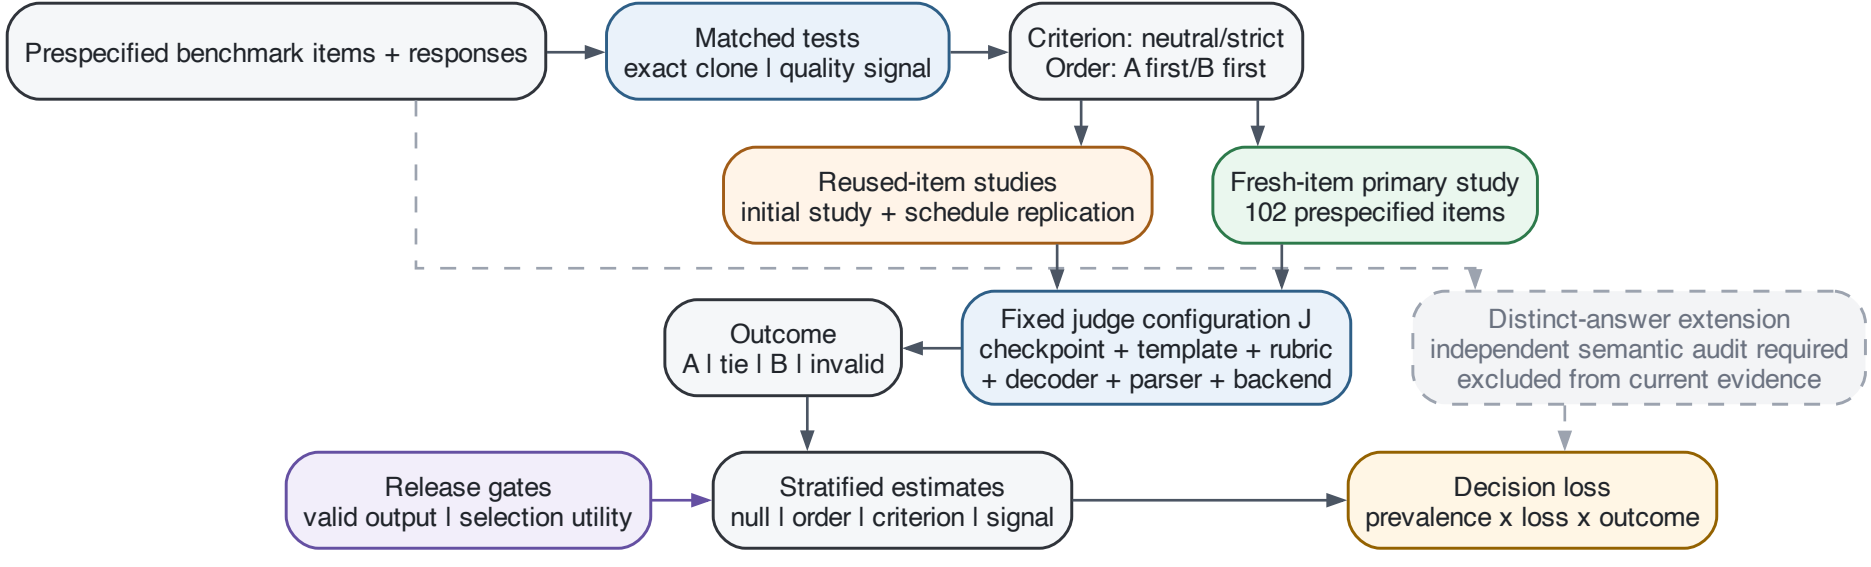
Prespecified benchmark items + responses
Matched tests
exact clone | quality signal
Criterion: neutral/strict
Order: A first/B first
Reused-item studies
initial study + schedule replication
Fresh-item primary study
102 prespecified items
Fixed judge configuration J
checkpoint + template + rubric
+ decoder + parser + backend
Outcome
A | tie | B | invalid
Distinct-answer extension
independent semantic audit required
excluded from current evidence
Release gates
valid output | selection utility
Stratified estimates
null | order | criterion | signal
Decision loss
prevalence x loss x outcome